



Deep Learning School

Физтех-Школа Прикладной математики и информатики (ФПМИ) МФТИ

Для быстрого выполнения просмотрите [семинар](#).

▼ Models: Sentence Sentiment Classification

Our goal is to create a model that takes a sentence (just like the ones in our dataset) and produces either 1 (indicating the sentence carries a positive sentiment) or a 0 (indicating the sentence carries a negative sentiment). More details of the task can be found [here](#).

Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)

“a visually stunning
rumination on love”



Movie Review
Sentiment Classifier



positive

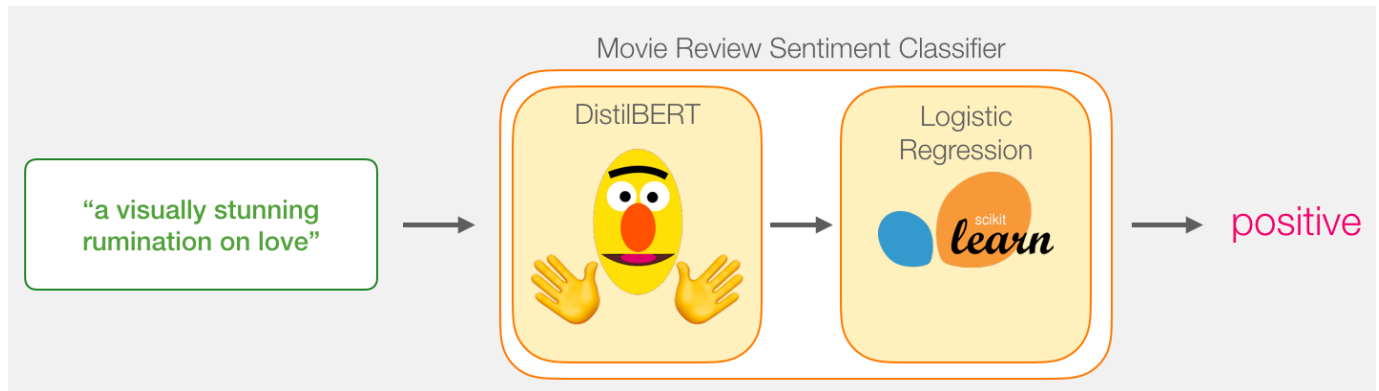
Under the hood, the model is actually made up of two model.

- DistilBERT processes the sentence and passes along some information it extracted from it on to the next model. DistilBERT is a smaller version of BERT developed and open sourced

by the team at HuggingFace. It's a lighter and faster version of BERT that roughly matches its performance.

- The next model, a basic Logistic Regression model from scikit learn will take in the result of DistilBERT's processing, and classify the sentence as either positive or negative (1 or 0, respectively).

The data we pass between the two models is a vector of size 768. We can think of this of vector as an embedding for the sentence that we can use for classification.



Dataset

The dataset we will use in this example is [SST2](#), which contains sentences from movie reviews, each labeled as either positive (has the value 1) or negative (has the value 0):

sentence	label
a stirring , funny and finally transporting re imagining of beauty and the beast and 1930s horror films	1
apparently reassembled from the cutting room floor of any given daytime soap	0
they presume their audience won't sit still for a sociology lesson	0
this is a visually stunning rumination on love , memory , history and the war between art and commerce	1
jonathan parker 's bartleby should have been the be all end all of the modern office anomie films	1

Installing the transformers library

Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)

ep learning

```
1 !pip install transformers
```

```
Collecting transformers
```

```
  Downloading transformers-4.12.5-py3-none-any.whl (3.1 MB)
```

```
    |████████████████████████████████████████| 3.1 MB 5.2 MB/s
```

```
Collecting sacremoses
```

```
  Downloading sacremoses-0.0.46-py3-none-any.whl (895 kB)
```

```
    |████████████████████████████████████████| 895 kB 27.0 MB/s
```

```
Requirement already satisfied: filelock in /usr/local/lib/python3.7/dist-packages (fr
```

```
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.7/dist-packages
```

```
Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.7/dist-pac
```

```
Collecting tokenizers<0.11,>=0.10.1
```

```
  Downloading tokenizers-0.10.3-cp37-cp37m-manylinux_2_5_x86_64.manylinux1_x86_64.mar
```

```
|████████████████████████████████████████| 3.3 MB 36.5 MB/s
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.7/dist-packages (from transformers==4.10.0)
Collecting huggingface-hub<1.0,>=0.1.0
  Downloading huggingface_hub-0.1.2-py3-none-any.whl (59 kB)
|████████████████████████████████████████| 59 kB 6.5 MB/s
Requirement already satisfied: importlib-metadata in /usr/local/lib/python3.7/dist-packages (from transformers==4.10.0)
Requirement already satisfied: tqdm>=4.27 in /usr/local/lib/python3.7/dist-packages (from transformers==4.10.0)
Collecting pyyaml>=5.1
  Downloading PyYAML-6.0-cp37-cp37m-manylinux_2_5_x86_64.manylinux1_x86_64.manylinux_2_17_x86_64.manylinux2014_x86_64.whl (596 kB)
|████████████████████████████████████████| 596 kB 39.9 MB/s
Requirement already satisfied: requests in /usr/local/lib/python3.7/dist-packages (from transformers==4.10.0)
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.7/dist-packages (from transformers==4.10.0)
Requirement already satisfied: pyparsing<3,>=2.0.2 in /usr/local/lib/python3.7/dist-packages (from transformers==4.10.0)
Requirement already satisfied: zipp>=0.5 in /usr/local/lib/python3.7/dist-packages (from transformers==4.10.0)
Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.7/dist-packages (from transformers==4.10.0)
Requirement already satisfied: urllib3!=1.25.0,!>=1.25.1,<1.26,>=1.21.1 in /usr/local/lib/python3.7/dist-packages (from transformers==4.10.0)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.7/dist-packages (from transformers==4.10.0)
Requirement already satisfied: chardet<4,>=3.0.2 in /usr/local/lib/python3.7/dist-packages (from transformers==4.10.0)
Requirement already satisfied: click in /usr/local/lib/python3.7/dist-packages (from transformers==4.10.0)
Requirement already satisfied: six in /usr/local/lib/python3.7/dist-packages (from transformers==4.10.0)
Requirement already satisfied: joblib in /usr/local/lib/python3.7/dist-packages (from transformers==4.10.0)
Installing collected packages: pyyaml, tokenizers, sacremoses, huggingface-hub, transformers
Attempting uninstall: pyyaml
  Found existing installation: PyYAML 3.13
  Uninstalling PyYAML-3.13:
    Successfully uninstalled PyYAML-3.13
Successfully installed huggingface-hub-0.1.2 pyyaml-6.0 sacremoses-0.0.46 tokenizers-0.10.2 transformers-4.10.0
```

[Transformers library doc](#)

Automatic saving failed. This file was updated remotely or in another tab.

[Show](#)

[diff](#)



HUGGING FACE

On a mission to solve NLP,
one commit at a time.



Star

36,299

```
1 import numpy as np
2 import pandas as pd
3
4 from sklearn.model_selection import train_test_split
5 from sklearn.linear_model import LogisticRegression
6 from sklearn.model_selection import GridSearchCV
7 from sklearn.model_selection import cross_val_score
```

Automatic saving failed. This file was updated remotely or in another tab.

[Show](#)

```
diff
10 from torch import nn
11 import torch.nn.functional as F
12
13 import transformers as ppb
14
15 import warnings
16 warnings.filterwarnings('ignore')
```

▼ Importing the dataset

```
1 df = pd.read_csv(
2     'https://github.com/clairett/pytorch-sentiment-classification/raw/master/data/SST2/
3     delimiter='\t',
```

```

4     header=None
5 )
6 print(df.shape)
7 df.head()

(6920, 2)


```

		0	1
0	a stirring , funny and finally transporting re...	1	
1	apparently reassembled from the cutting room f...	0	
2	they presume their audience wo n't sit still f...	0	
3	this is a visually stunning rumination on love...	1	
4	jonathan parker 's bartleby should have been t...	1	

▼ Using BERT for text classification.

Let's now load a pre-trained BERT model.

```

1 # For DistilBERT, Load pretrained model/tokenizer:
2
3 model_class, tokenizer_class, pretrained_weights = (ppb.DistilBertModel, ppb.DistilBert
4 tokenizer = tokenizer_class.from_pretrained(pretrained_weights)
5 model = model_class.from_pretrained(pretrained_weights)

```

```

Downloading: 100% 226k/226k [00:00<00:00, 754kB/s]
Downloading: 100% 28.0/28.0 [00:00<00:00, 505B/s]
Downloading: 100% 455k/455k [00:00<00:00, 1.02MB/s]
Downloading: 100% 483/483 [00:00<00:00, 9.00kB/s]
Downloading: 100% 256M/256M [00:09<00:00, 26.9MB/s]

```

Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)

it used when ir
ckpoint of a n
checkpoint of

```

1 # look at the model
2 device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
3 model = model.to(device)
4 model.eval()

      (lin2): Linear(in_features=3072, out_features=768, bias=True)
    )
  (output_layer_norm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
)
(2): TransformerBlock(
  (attention): MultiHeadSelfAttention(
    (dropout): Dropout(p=0.1, inplace=False)
    (q_lin): Linear(in_features=768, out_features=768, bias=True)
    (k_lin): Linear(in_features=768, out_features=768, bias=True)

```

```

        (k_lin): Linear(in_features=768, out_features=768, bias=True)
        (v_lin): Linear(in_features=768, out_features=768, bias=True)
        (out_lin): Linear(in_features=768, out_features=768, bias=True)
    )
    (sa_layer_norm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
    (ffn): FFN(
        (dropout): Dropout(p=0.1, inplace=False)
        (lin1): Linear(in_features=768, out_features=3072, bias=True)
        (lin2): Linear(in_features=3072, out_features=768, bias=True)
    )
    (output_layer_norm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
)
(3): TransformerBlock(
  (attention): MultiHeadSelfAttention(
    (dropout): Dropout(p=0.1, inplace=False)
    (q_lin): Linear(in_features=768, out_features=768, bias=True)
    (k_lin): Linear(in_features=768, out_features=768, bias=True)
    (v_lin): Linear(in_features=768, out_features=768, bias=True)
    (out_lin): Linear(in_features=768, out_features=768, bias=True)
  )
  (sa_layer_norm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
  (ffn): FFN(
    (dropout): Dropout(p=0.1, inplace=False)
    (lin1): Linear(in_features=768, out_features=3072, bias=True)
    (lin2): Linear(in_features=3072, out_features=768, bias=True)
  )
  (output_layer_norm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
)
(4): TransformerBlock(
  (attention): MultiHeadSelfAttention(
    (dropout): Dropout(p=0.1, inplace=False)
    (q_lin): Linear(in_features=768, out_features=768, bias=True)
    (k_lin): Linear(in_features=768, out_features=768, bias=True)
    (v_lin): Linear(in_features=768, out_features=768, bias=True)
    (out_lin): Linear(in_features=768, out_features=768, bias=True)
  )
  (sa_layer_norm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
  (ffn): FFN(
    (dropout): Dropout(p=0.1, inplace=False)
    (lin1): Linear(in_features=768, out_features=3072, bias=True)
    (lin2): Linear(in_features=3072, out_features=768, bias=True)
  )
  (output_layer_norm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
)
(5): TransformerBlock(
  (attention): MultiHeadSelfAttention(
    (dropout): Dropout(p=0.1, inplace=False)
    (q_lin): Linear(in_features=768, out_features=768, bias=True)
    (k_lin): Linear(in_features=768, out_features=768, bias=True)
    (v_lin): Linear(in_features=768, out_features=768, bias=True)
    (out_lin): Linear(in_features=768, out_features=768, bias=True)
  )
  (sa_layer_norm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
  (ffn): FFN(
    (dropout): Dropout(p=0.1, inplace=False)
    (lin1): Linear(in_features=768, out_features=3072, bias=True)
    (lin2): Linear(in_features=3072, out_features=768, bias=True)
  )
  (output_layer_norm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
)

```

ffine=True)

Automatic saving failed. This file was updated remotely or in another tab.

[Show](#)

[diff](#)

```

1 from termcolor import colored
2
3 colors = ['red', 'green', 'blue', 'yellow']
4
5 def model_structure(layer, margin=0, item_color=0):
6     for name, next_layer in layer.named_children():
7

```

```

8         next = (0 if not list(next_layer.named_children()) else 1)
9         print(colored(' ' * margin + name, colors[item_color]) + ':' * next)
10        model_structure(next_layer, margin + len(name) + 2, (item_color + 1) % 4)
11
12 model_structure(model)

```

```

1:
    attention:
        dropout
        q_lin
        k_lin
        v_lin
        out_lin
    sa_layer_norm
    ffn:
        dropout
        lin1
        lin2
    output_layer_norm
2:
    attention:
        dropout
        q_lin
        k_lin
        v_lin
        out_lin
    sa_layer_norm
    ffn:
        dropout
        lin1
        lin2
    output_layer_norm
3:
    attention:
        dropout
        q_lin
        k_lin
        v_lin
        out_lin
    sa_layer_norm
    ffn:

```

Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)

```

    output_layer_norm
4:
    attention:
        dropout
        q_lin
        k_lin
        v_lin
        out_lin
    sa_layer_norm
    ffn:
        dropout
        lin1
        lin2
    output_layer_norm
5:
    attention:
        dropout

```

```
q_lin
k_lin
v_lin
out_lin
```

▼ Preparing the dataset

```
1 from torch.utils.data import Dataset, random_split
2
3 class ReviewsDataset(Dataset):
4     def __init__(self, reviews, tokenizer, labels):
5         self.labels = labels
6         # tokenized reviews
7         self.tokenized = reviews.apply((lambda x: tokenizer.encode(x, add_special_token
8
9     def __getitem__(self, idx):
10         return {"tokenized": self.tokenized[idx], "label": self.labels[idx]}
11
12     def __len__(self):
13         return len(self.labels)
14
15 dataset = ReviewsDataset(reviews=df[0],
16                           labels=df[1],
17                           tokenizer=tokenizer)
18
19 # DON'T CHANGE, PLEASE
20 train_size, val_size = int(.8 * len(dataset)), int(.1 * len(dataset))
21 torch.manual_seed(2)
22 train_data, valid_data, test_data = random_split(dataset, [train_size, val_size, len(da
23
24 print(f"Number of training examples: {len(train_data)}")
25 print(f"Number of validation examples: {len(valid_data)}")
26 print(f"Number of testing examples: {len(test_data)}")
```

```
Number of training examples: 5536
Number of validation examples: 692
```

Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)

```
1 from torch.utils.data import Sampler
2
3 class ReviewsSampler(Sampler):
4     def __init__(self, subset, batch_size=32):
5         self.batch_size = batch_size
6         self.subset = subset
7
8         self.indices = subset.indices
9         # tokenized for our data
10        self.tokenized = np.array(subset.dataset.tokenized)[self.indices]
11
12    def __iter__(self):
13
14        batch_idx = []
```



```

15         # index in sorted data
16         for index in np.argsort(list(map(len, self.tokenized))):
17             batch_idx.append(index)
18             if len(batch_idx) == self.batch_size:
19                 yield batch_idx
20                 batch_idx = []
21
22         if len(batch_idx) > 0:
23             yield batch_idx
24
25     def __len__(self):
26         return len(self.dataset)

1 from torch.utils.data import DataLoader
2
3 def get_padded(values):
4     max_len = 0
5     for value in values:
6         if len(value) > max_len:
7             max_len = len(value)
8
9     padded = np.array([value + [0]*(max_len-len(value)) for value in values])
10
11     return padded
12
13 def collate_fn(batch):
14
15     inputs = []
16     labels = []
17     for elem in batch:
18         inputs.append(elem['tokenized'])
19         labels.append(elem['label'])
20
21     inputs = get_padded(inputs) # padded inputs
22     attention_mask = np.where(inputs != 0, 1, 0)
23
24     return {"inputs": torch.tensor(inputs),

```

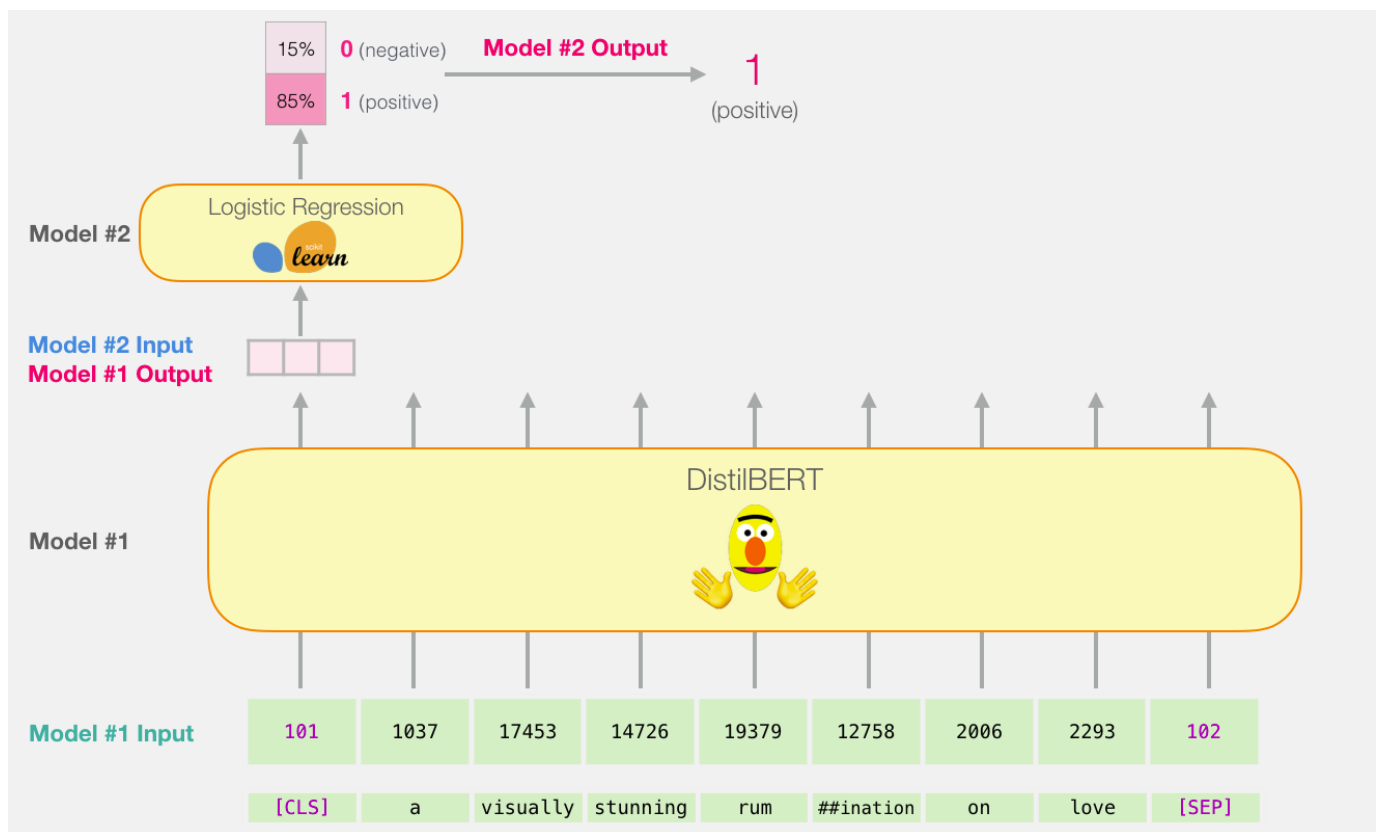
Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)

```

28 train_loader = DataLoader(train_data, batch_sampler=ReviewsSampler(train_data), collate
29 valid_loader = DataLoader(valid_data, batch_sampler=ReviewsSampler(valid_data), collate
30 test_loader = DataLoader(test_data, batch_sampler=ReviewsSampler(test_data), collate_fn

```

▼ Baseline



```

1 from tqdm.notebook import tqdm
2
3 def get_xy(loader):
4     features = []
5     labels = []
6
7     with torch.no_grad():
8         for batch in tqdm(loader):
9
10            # don't forget about .to(device)
11            batch_inputs = batch['inputs'].to(device)
12            batch_attention_mask = batch['attention_mask'].to(device)
13            # no need to put on GPU for now
14
15            # forward
16            last_hidden_states = model(input_ids=batch_inputs,
17                                     attention_mask=batch_attention_mask)
18
19            # append features and labels
20            features.append(last_hidden_states[0].cpu())
21            labels.append(batch_labels)
22
23
24    features = torch.cat([elem[:, 0, :] for elem in features], dim=0).numpy()
25    labels = torch.cat(labels, dim=0).numpy()
26
27    return features, labels

```

Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)

```

1 train_features, train_labels = get_xy(train_loader)

```

```

2 valid_features, valid_labels = get_xy(valid_loader)
3 test_features, test_labels = get_xy(test_loader)

```

```
173/? [00:08<00:00, 11.52it/s]
```

```
22/? [00:01<00:00, 16.76it/s]
```

```
22/? [00:01<00:00, 15.85it/s]
```

```

1 lr_clf = LogisticRegression()
2 lr_clf.fit(train_features, train_labels)
3 lr_clf.score(test_features, test_labels)

```

```
0.815028901734104
```

▼ Fine-Tuning BERT

Define the model

```

1 from torch import nn
2
3 class BertClassifier(nn.Module):
4     def __init__(self, pretrained_model, dropout=0.1):
5         super().__init__()
6
7         self.bert = pretrained_model
8         self.dropout = nn.Dropout(p=dropout)
9         self.relu = nn.ReLU()
10
11        self.out = nn.Linear(in_features=768, out_features=1)
12
13    def forward(self, inputs, attention_mask):
14        last_hidden_states = model(input_ids=inputs, attention_mask=attention_mask)[0]
15        # [batch size, max seq len, bert hidden dim]
16
17        features = last_hidden_states[:, 0, :] # [batch size, bert hidden dim]
18        proba = self.out(self.relu(self.dropout(features))) # [batch size, 1]
19
20        proba = proba.view(-1) # [batch_size, ] - probability to be positive
21
22        return proba

```

Automatic saving failed. This file was updated remotely or in another tab.

[Show](#)

[diff](#)

```

1 import torch.optim as optim
2
3 # DON'T CHANGE
4 model = model_class.from_pretrained(pretrained_weights).to(device)
5 bert_clf = BertClassifier(model).to(device)
6 # you can change
7 optimizer = optim.Adam(bert_clf.parameters(), lr=2e-5)
8 criterion = nn.BCEWithLogitsLoss()

```

Some weights of the model checkpoint at distilbert-base-uncased were not used when in

- This IS expected if you are initializing DistilBertModel from the checkpoint of a n
- This IS NOT expected if you are initializing DistilBertModel from the checkpoint of

```
1 def train(model, iterator, optimizer, criterion, clip, train_history=None, valid_history=None):
2     model.train()
3
4     epoch_loss = 0
5     history = []
6     for i, batch in enumerate(iterator):
7
8         # don't forget about .to(device)
9         inputs = batch['inputs'].to(device)
10        labels = batch['labels'].to(device)
11        attention_mask = batch['attention_mask'].to(device)
12
13
14        optimizer.zero_grad()
15
16        # forward + backward + optimize
17        output = model(inputs, attention_mask)
18        loss = criterion(output, labels)
19        loss.backward()
20        torch.nn.utils.clip_grad_norm_(model.parameters(), clip)
21        optimizer.step()
22
23        epoch_loss += loss.item()
24
25        history.append(loss.cpu().data.numpy())
26        if (i+1)%10==0:
27            fig, ax = plt.subplots(nrows=1, ncols=2, figsize=(12, 8))
28
29            clear_output(True)
30            ax[0].plot(history, label='train loss')
31            ax[0].set_xlabel('Batch')
32            ax[0].set_title('Train loss')
33            if train_history is not None:
34                if valid_history is not None:
35                    ax[1].plot(valid_history, label='general valid history')
36                    plt.legend()
37
38            plt.show()
39
40            # # for debugging
41            # break
42
43
44
45
46    return epoch_loss / (i + 1)
47
48 def evaluate(model, iterator, criterion):
49
50    model.eval()
```

Automatic saving failed. This file was updated remotely or in another tab. [Show](#)

[diff](#)

```

51
52     epoch_loss = 0
53
54     history = []
55
56     with torch.no_grad():
57
58         for i, batch in enumerate(iterator):
59
60             inputs = batch['inputs'].to(device)
61             labels = batch['labels'].to(device)
62             attention_mask = batch['attention_mask'].to(device)
63
64             output = model(inputs, attention_mask)
65             loss = criterion(output, labels)
66
67             epoch_loss += loss.item()
68
69             # # for debugging
70             # break
71
72     return epoch_loss / (i + 1)
73
74 def epoch_time(start_time, end_time):
75     elapsed_time = end_time - start_time
76     elapsed_mins = int(elapsed_time / 60)
77     elapsed_secs = int(elapsed_time - (elapsed_mins * 60))
78     return elapsed_mins, elapsed_secs

```

```

1 import time
2 import math
3 import matplotlib
4 matplotlib.rcParams.update({'figure.figsize': (16, 12), 'font.size': 14})
5 import matplotlib.pyplot as plt
6 %matplotlib inline
7 from IPython.display import clear_output

```

Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)

```


3
4 N_EPOCHS = 3
5 CLIP = 1
6
7 best_valid_loss = float('inf')
8
9 for epoch in range(N_EPOCHS):
10
11     start_time = time.time()
12
13     train_loss = train(bert_clf, train_loader, optimizer, criterion, CLIP, train_history)
14     valid_loss = evaluate(bert_clf, valid_loader, criterion)
15
16     end_time = time.time()
17

```

```

18 epoch_mins, epoch_secs = epoch_time(start_time, end_time)
19
20 if valid_loss < best_valid_loss:
21     best_valid_loss = valid_loss
22     torch.save(bert_clf.state_dict(), 'best-val-model.pt')
23
24 train_history.append(train_loss)
25 valid_history.append(valid_loss)
26 print(f'Epoch: {epoch+1:02} | Time: {epoch_mins}m {epoch_secs}s')
27 print(f'\tTrain Loss: {train_loss:.3f} | Train PPL: {math.exp(train_loss):7.3f}')
28 print(f'\tVal. Loss: {valid_loss:.3f} | Val. PPL: {math.exp(valid_loss):7.3f}')

```



Epoch: 03 | Time: 0m 38s

Train Loss: 0.136 | Train PPL: 1.145

Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)

```

1 best_model = BertClassifier(model).to(device)
2 best_model.load_state_dict(torch.load('best-val-model.pt'))
3
4 pred_labels = []
5 true_labels = []
6
7 best_model.eval()
8 with torch.no_grad():
9     for i, batch in tqdm(enumerate(test_loader)):
10
11         inputs = batch['inputs'].to(device)
12         attention_mask = batch['attention_mask'].to(device)

```

```

13     labels = batch['labels'] # used with sklearn later => stays on cpu
14
15     output = best_model(inputs, attention_mask)
16     binary_output = (torch.sigmoid(output) > 0.5).to(torch.int)
17
18     true_labels.append(labels.numpy())
19     pred_labels.append(binary_output.cpu().numpy())

```

22/? [00:01<00:00, 13.91it/s]

```

1 from sklearn.metrics import accuracy_score
2
3 true_labels = np.concatenate(true_labels, axis=0)
4 pred_labels = np.concatenate(pred_labels, axis=0)
5 accuracy_score(true_labels, pred_labels)

```

```

-----
ValueError                                Traceback (most recent call last)
<ipython-input-21-6f540048575b> in <module>()
      1 from sklearn.metrics import accuracy_score
      2
----> 3 true_labels = np.concatenate(true_labels, axis=0)
      4 pred_labels = np.concatenate(pred_labels, axis=0)
      5 accuracy_score(true_labels, pred_labels)

<__array_function__ internals> in concatenate(*args, **kwargs)

ValueError: zero-dimensional arrays cannot be concatenated

```

SEARCH STACK OVERFLOW

```
1 assert accuracy_score(true_labels, pred_labels) >= 0.86
```

▼ Finetuned model from HUGGING FACE

BertForSequenceClassification

Automatic saving failed. This file was updated remotely or in another tab. [Show](#)

diff

```

1 from transformers import AutoTokenizer, AutoModelForSequenceClassification
2
3 # we have the same tokenizer
4 # new_tokenizer = AutoTokenizer.from_pretrained("distilbert-base-uncased-finetuned-sst-
5 new_model = AutoModelForSequenceClassification.from_pretrained("distilbert-base-uncased

```

Downloading: 100%

629/629 [00:00<00:00, 16.1kB/s]

Downloading: 100%

255M/255M [00:09<00:00, 31.2MB/s]

```


1 pred_labels = []
2 true_labels = []
3
4 new_model.eval()
5 with torch.no_grad():

```

```

5 with torch.no_grad():
6     for i, batch in tqdm(enumerate(test_loader)):
7
8         inputs = batch['inputs'].to(device)
9         attention_mask = batch['attention_mask'].to(device)
10        labels = batch['labels'] # used with sklearn later => stays on cpu
11
12        output = best_model(inputs, attention_mask)
13        binary_output = (torch.sigmoid(output) > 0.5).to(torch.int)
14
15        true_labels.append(labels.numpy())
16        pred_labels.append(binary_output.cpu().numpy())
17
18 true_labels = np.concatenate(true_labels, axis=0)
19 pred_labels = np.concatenate(pred_labels, axis=0)
20 accuracy_score(true_labels, pred_labels)

```

 22/? [00:01<00:00, 13.60it/s]
0.869942196531792

```

1 model_structure(new_model)

```

```

distilbert:
    embeddings:
        word_embeddings
        position_embeddings
        LayerNorm
        dropout
    transformer:
        layer:
            0:
                attention:
                    dropout
                    q_lin
                    k_lin
                    v_lin
                    out_lin
                sa_layer_norm
                ffn:

```

Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)

```

                    lin2
                output_layer_norm
            1:
                attention:
                    dropout
                    q_lin
                    k_lin
                    v_lin
                    out_lin
                sa_layer_norm
                ffn:
                    dropout
                    lin1
                    lin2
                output_layer_norm
            2:
                attention:

```



```

        dropout
        q_lin
        k_lin
        v_lin
        out_lin
    sa_layer_norm
    ffn:
        dropout
        lin1
        lin2
    output_layer_norm
3:
    attention:
        dropout
        q_lin
        k_lin
        v_lin
        out_lin
    sa_layer_norm
    ffn:
        dropout
        lin1
        lin2

```

Напишите вывод о своих результатах. В выводы включите ваши гиперпараметры.

Качество с помощью *Fine-Tuning* должно достигать 0.86. Все происходит крайне быстро, в отличие от gpt. Качество предобученной модели и baseline не сильно отличается.

1

Automatic saving failed. This file was updated remotely or in another tab.

[diff](#)

[Show](#)

